

OpenACC

Open Accelerators

Matthew Silva
Professor Alvarez
CSC 415
9 April 2018

“What is OpenACC?”

OpenACC is a user-driven **directive-based performance-portable parallel programming model** designed for scientists and engineers interested in **porting their codes** to a wide-variety of **heterogeneous [CPU/GPU] HPC hardware** platforms and architectures with **significantly less programming effort than required with a low-level model.**”

Who remembers the programming effort required for other low-level tasks?

(AVX isn't even very low-level.)

Source: “What Is OpenACC?”

Heterogeneous Computing Review

“...**conventional** and **specialized** processors work cooperatively.”

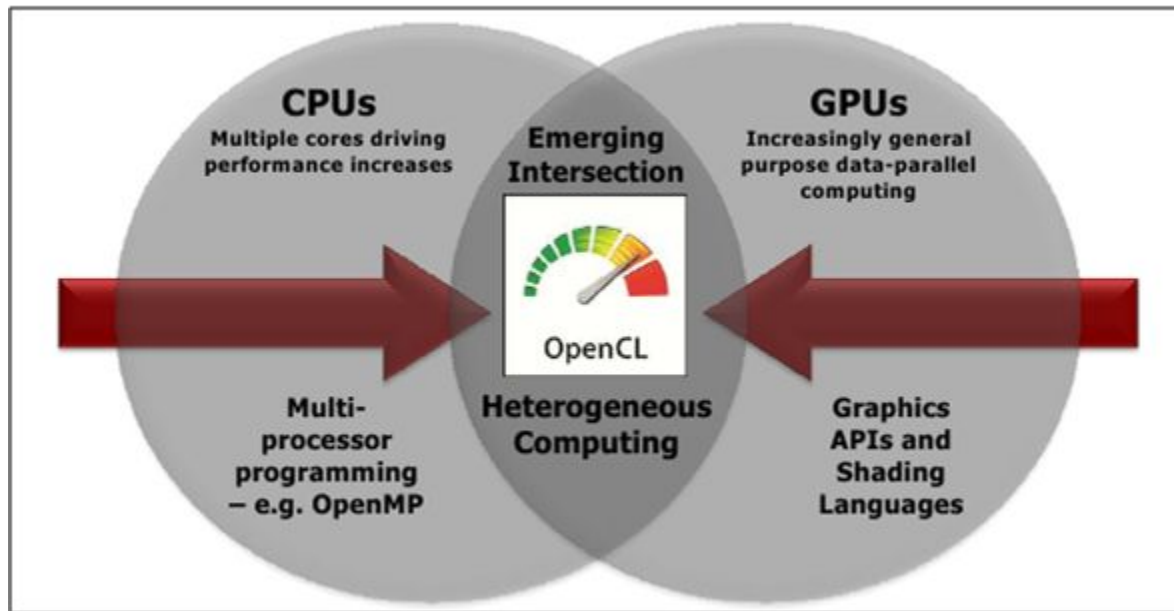
“Conventional processors” generally refers to a regular CPU.

“Specialized processors” may refer to a GPU or other processor meant to perform a certain type of task.

Heterogeneous computing allows conventional and specialized processors to team up, each solving the parts of a problem that they are best at.

OpenACC will accelerate tasks by assigning them to the specialized processor.

Heterogeneous Computing Review



OpenCL vs OpenACC

OpenCL and OpenACC both deal with parallelizing code for heterogeneous computing applications.

OpenACC is much easier to use, since it relies almost fully on compiler directives that find and create parallelism for you.

(Much less work, but also less effective.)

You can fine-tune OpenACC directives, but many feel that this removes the abstraction that creates appeal for OpenACC over OpenCL.

Source: Melonakos

Directives (pragma)

```
#pragma acc data copy(A) create(Anew)
while ( error > tol  &&  iter <  iter_max ) {
    error = 0.0;
    #pragma acc kernels {
    #pragma acc loop independent collapse(2)
        for (  int  j = 1; j < n-1; j++ ) {
            for (  int  i = 1; i < m-1; i++ ) {
                Anew [j] [i] = 0.25 * ( A [j] [i+1] + A [j] [i-1] +
                                         A [j-1] [i] + A [j+1] [i]);
                error = max ( error, fabs (Anew [j] [i] - A [j] [i]));
            }
        }
    }
}
```

Directives Review

Directives are comments that act as commands to the compiler.

Directives are NOT code! (At least, not code in the language being compiled.)

Directives tell the compiler to try to compile the code a certain way (make it faster.)

Directive Results (Compiler Output)

```
20, Generated vector simd code for the loop containing reductions
Generated 2 prefetch instructions for the loop
FMA (fused multiply-add) instruction(s) generated

matvec(const float *, const float *, unsigned int, unsigned int, float *):
20, Accelerator kernel generated
    Generating Tesla code
        22, #pragma acc loop gang, vector(128) /* blockIdx.x threadIdx.x */
        24, Generating implicit reduction(+:vecSum)
20, Generating implicit copyin(A[in:n],x[:n])
```

OpenACC Directives are a portable way to show OpenACC where to automatically generate vectorized code for a specified accelerator

Here, tesla was selected while compiling.

Directive Results (Compiler Output)

```
20, Generated vector simd code for the loop containing reductions
Generated 2 prefetch instructions for the loop
FMA (fused multiply-add) instruction(s) generated

matvec(const float *, const float *, unsigned int, unsigned int, float *):
20, Accelerator kernel generated
   Generating Tesla code
   22, #pragma acc loop gang, vector(128) /* blockIdx.x threadIdx.x */
   24, Generating implicit reduction(+:vecSum)
20, Generating implicit copyin(A[in:n],x[:n])
```

Notice accelerator-specific parallelism being created with SIMD. The tesla is able to use SIMD.

An OpenACC gang is a thread block organized to work together on a task, commonly used for GPUs and other manycore accelerators.

The copyin is used to reduce the disadvantage of accelerator memory latency with the CPU (Move all the data to the GPU at once).

Runtime Functions

OpenACC also provides some runtime functions (like AVX and AVX2).

Examples:

- `acc_alloc()`
- `acc_free()`

Source: “OpenACC API”

Who can use OpenACC?

“The OpenACC specification supports **C, C++, Fortran** programming languages and multiple hardware architectures including **X86 & POWER CPUs, NVIDIA GPUs, and Xeon KNL** in the near future.”

Many compilers (both free and commercial) now support OpenACC directives.

Sources: “Get Started” and “Downloads & Tools”

Who can use OpenACC?

OpenACC is generally used for scientific applications.

“...take advantage of compilers and libraries to quickly accelerate your codes with CPUs and GPUs so that you can spend more time on real breakthroughs.”

Directives are an abstraction that removes much of the need for knowledge of hardware acceleration.

Scientists get the benefit of acceleration without having to change their programs or know about parallelism/hardware.

How to use OpenACC

OpenACC.org breaks up usage into three steps that can make optimizing a program for heterogeneous computing systems a simple task:

1. Analyze
2. Parallelize
3. Optimize

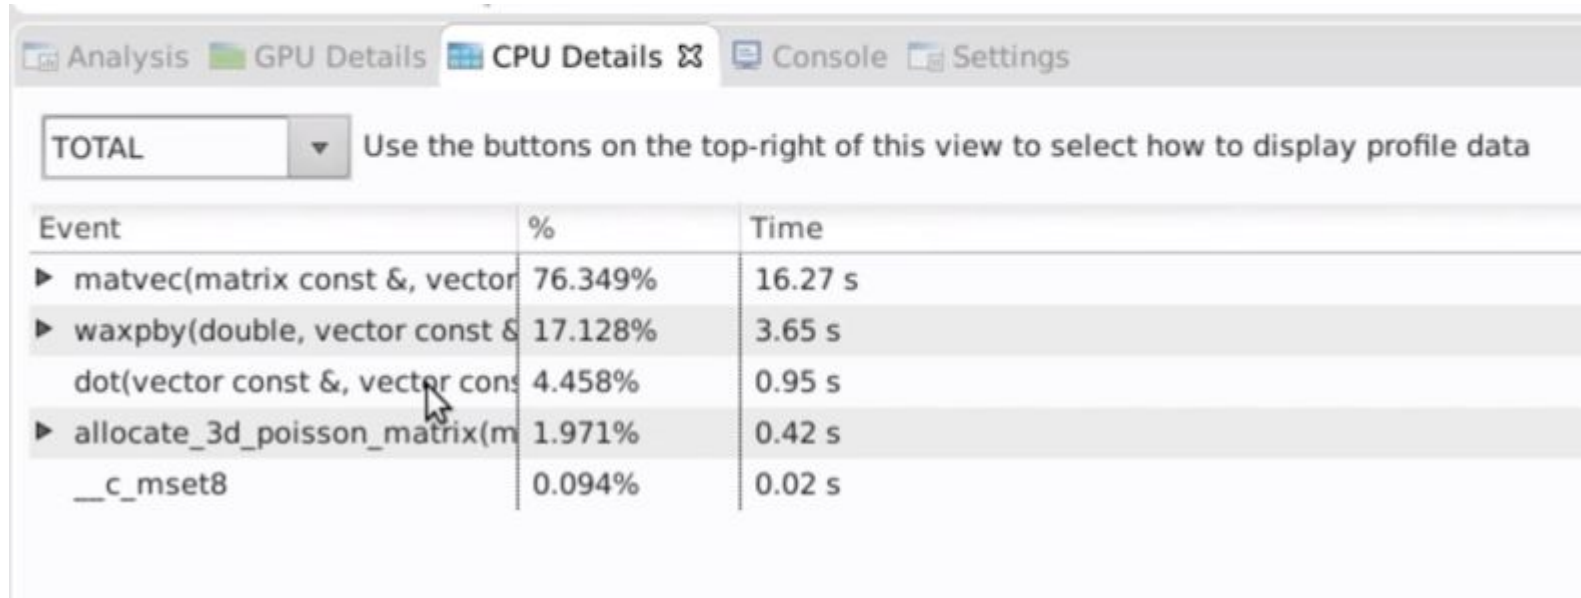
Analyze

OpenACC.org recommends using profiling tools to figure out where a program could benefit from acceleration.

Profiling tools run an executable, recording how much execution time is spent in each function.

Profiling tools can also compare the executable to the source code to see where a compiler made optimizations.

Analyze



Analysis GPU Details CPU Details Console Settings

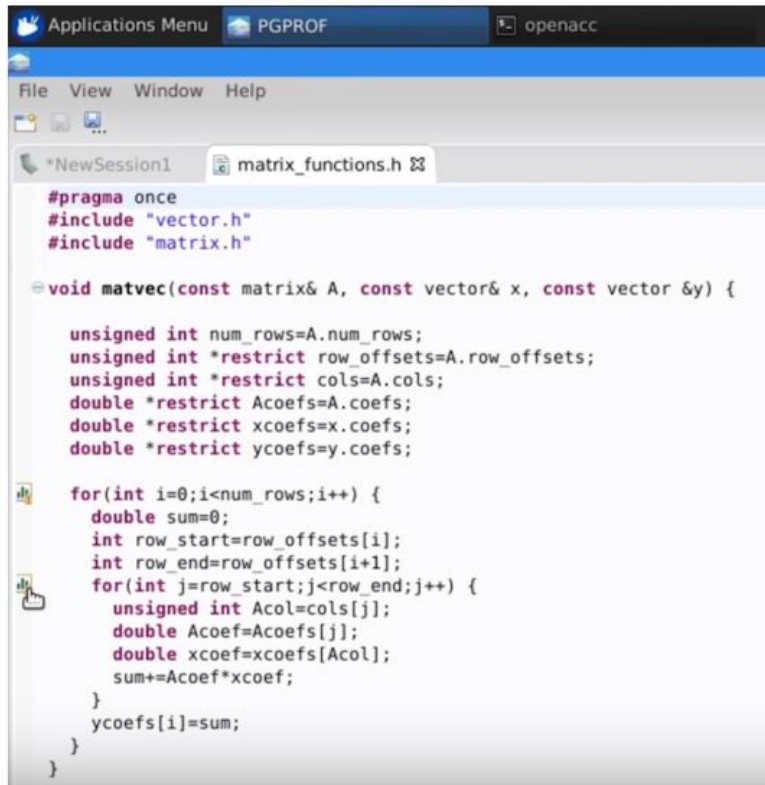
TOTAL Use the buttons on the top-right of this view to select how to display profile data

Event	%	Time
▶ matvec(matrix const &, vector	76.349%	16.27 s
▶ waxpby(double, vector const &	17.128%	3.65 s
dot(vector const &, vector cons	4.458%	0.95 s
▶ allocate_3d_poisson_matrix(m	1.971%	0.42 s
__c_mset8	0.094%	0.02 s

“...by accelerating the three most time consuming routines, I’ll accelerate nearly 98% of the application.”

Source: “Step 1: Analyze” (2:22)

Analyze

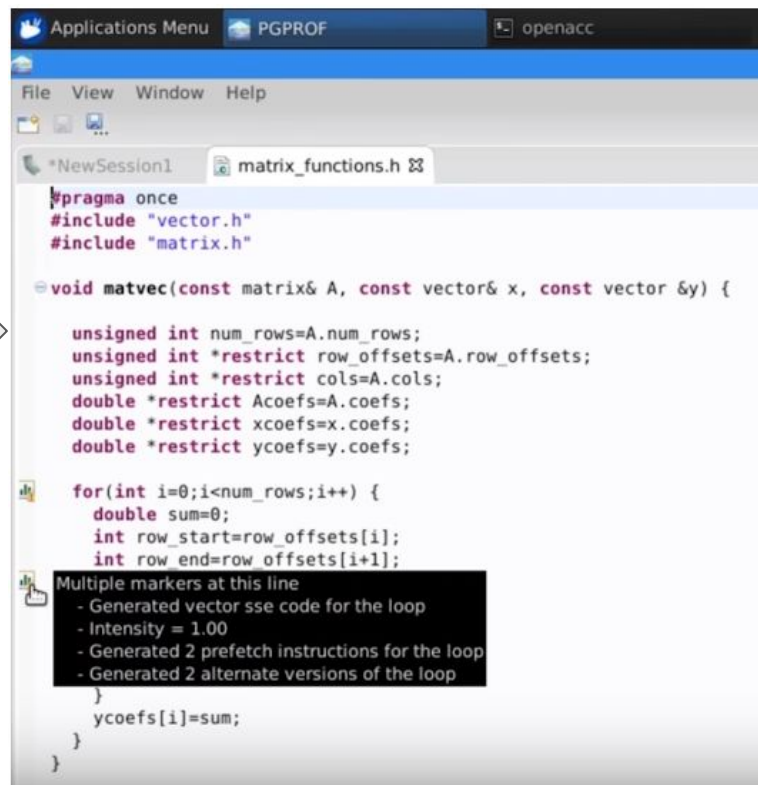



```
#pragma once
#include "vector.h"
#include "matrix.h"

void matvec(const matrix& A, const vector& x, const vector &y) {

    unsigned int num_rows=A.num_rows;
    unsigned int *restrict row_offsets=A.row_offsets;
    unsigned int *restrict cols=A.cols;
    double *restrict Acoefs=A.coefs;
    double *restrict xcoefs=x.coefs;
    double *restrict ycoefs=y.coefs;

    for(int i=0;i<num_rows;i++) {
        double sum=0;
        int row_start=row_offsets[i];
        int row_end=row_offsets[i+1];
        for(int j=row_start;j<row_end;j++) {
            unsigned int Acol=cols[j];
            double Acoef=Acoefs[j];
            double xcoef=xcoefs[Acol];
            sum+=Acoef*xcoef;
        }
        ycoefs[i]=sum;
    }
}
```



```
#pragma once
#include "vector.h"
#include "matrix.h"

void matvec(const matrix& A, const vector& x, const vector &y) {

    unsigned int num_rows=A.num_rows;
    unsigned int *restrict row_offsets=A.row_offsets;
    unsigned int *restrict cols=A.cols;
    double *restrict Acoefs=A.coefs;
    double *restrict xcoefs=x.coefs;
    double *restrict ycoefs=y.coefs;

    for(int i=0;i<num_rows;i++) {
        double sum=0;
        int row_start=row_offsets[i];
        int row_end=row_offsets[i+1];
        for(int j=row_start;j<row_end;j++) {
            unsigned int Acol=cols[j];
            double Acoef=Acoefs[j];
            double xcoef=xcoefs[Acol];
            sum+=Acoef*xcoef;
        }
        ycoefs[i]=sum;
    }
}
```

Multiple markers at this line

- Generated vector sse code for the loop
- Intensity = 1.00
- Generated 2 prefetch instructions for the loop
- Generated 2 alternate versions of the loop

The profiler examines the source code to show where and how the compiler performed optimizations.

Source: "Step 1: Analyze" (2:37)

Parallelize

Knowing where the compiler made optimizations from the profiler, a researcher has an idea of where the most work is being done, as well as loops/operations that can likely be parallelized.

With OpenACC, accelerating the program is now as easy as adding the proper compiler directives.

Parallelize

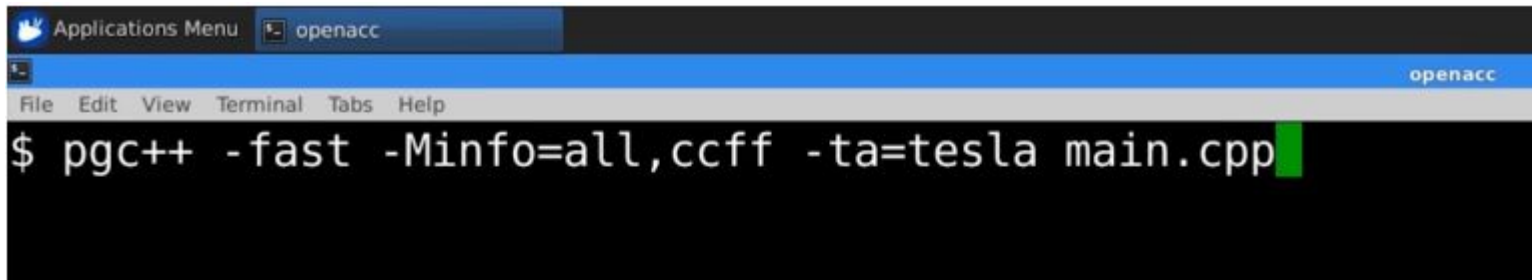
```
    sum += xcoefs[i] * ycoefs[i];  
}  
return sum;  
}  
  
void waxyby(double alpha, const vector &x, double beta, const vector &y, vector &w)  
{  
    unsigned int n=x.n;  
    double *restrict xcoefs=x.coefs;  
    double *restrict ycoefs=y.coefs;  
    double *restrict wcoefs=w.coefs;  
  
    #pragma acc parallel loop  
    for(int i=0;i<n;i++) {  
        wcoefs[i]=alpha*xcoefs[i]+beta*ycoefs[i];  
    }  
}
```

This pragma will accelerate the loop below it, automatically creating parallelism. No need to code unrolling or separate accumulators! Compiler does it all for you!

Parallelize

Very important to note that OpenACC will automatically tailor its parallelization to the target hardware.

No knowledge of your hardware is needed besides what kind of hardware it is.

A screenshot of a terminal window titled 'openacc'. The window has a menu bar with 'File', 'Edit', 'View', 'Terminal', 'Tabs', and 'Help'. The terminal text shows a command being entered: '\$ pgc++ -fast -Minfo=all,ccff -ta=tesla main.cpp'. A green cursor is at the end of the command.

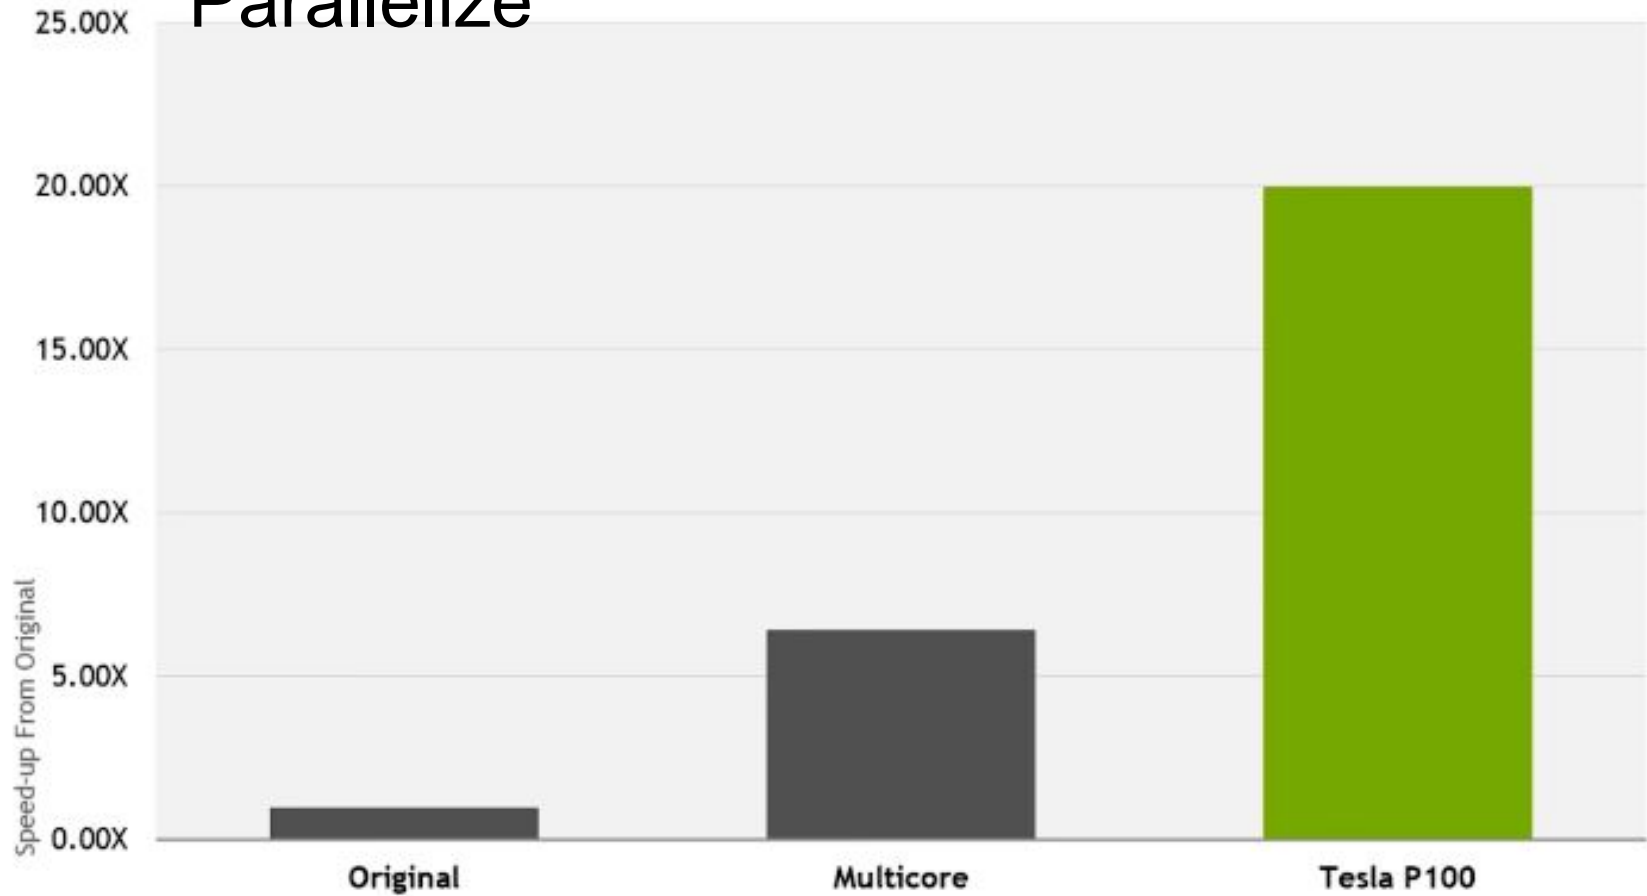
```
Applications Menu openacc  
openacc  
File Edit View Terminal Tabs Help  
$ pgc++ -fast -Minfo=all,ccff -ta=tesla main.cpp
```

All the researcher has to do is specify their type of specialized processor (target accelerator) with `-ta=<processor>` when compiling their code.

Parallelize

The compiler will follow the given directives to produce code that is parallelized where requested for optimal use by the target accelerator when possible.

Parallelize

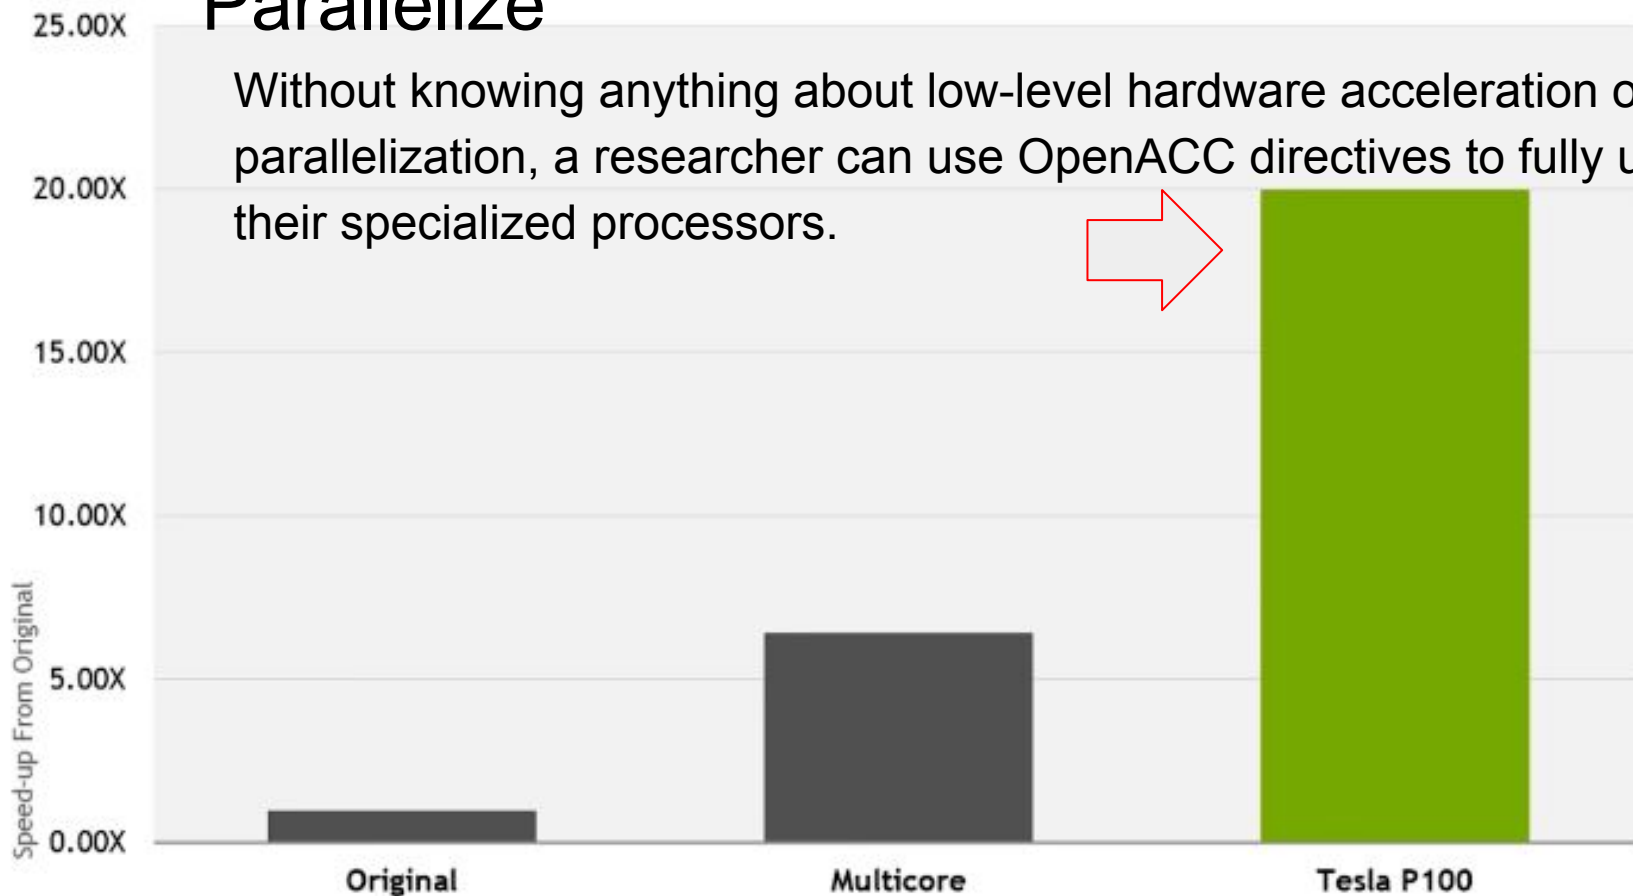


Source: PGI 16.10, Original and Multicore: Intel Xeon CPU E5-2698 v3 @ 2.30GHz

Source: "Step 2: Parallelize"

Parallelize

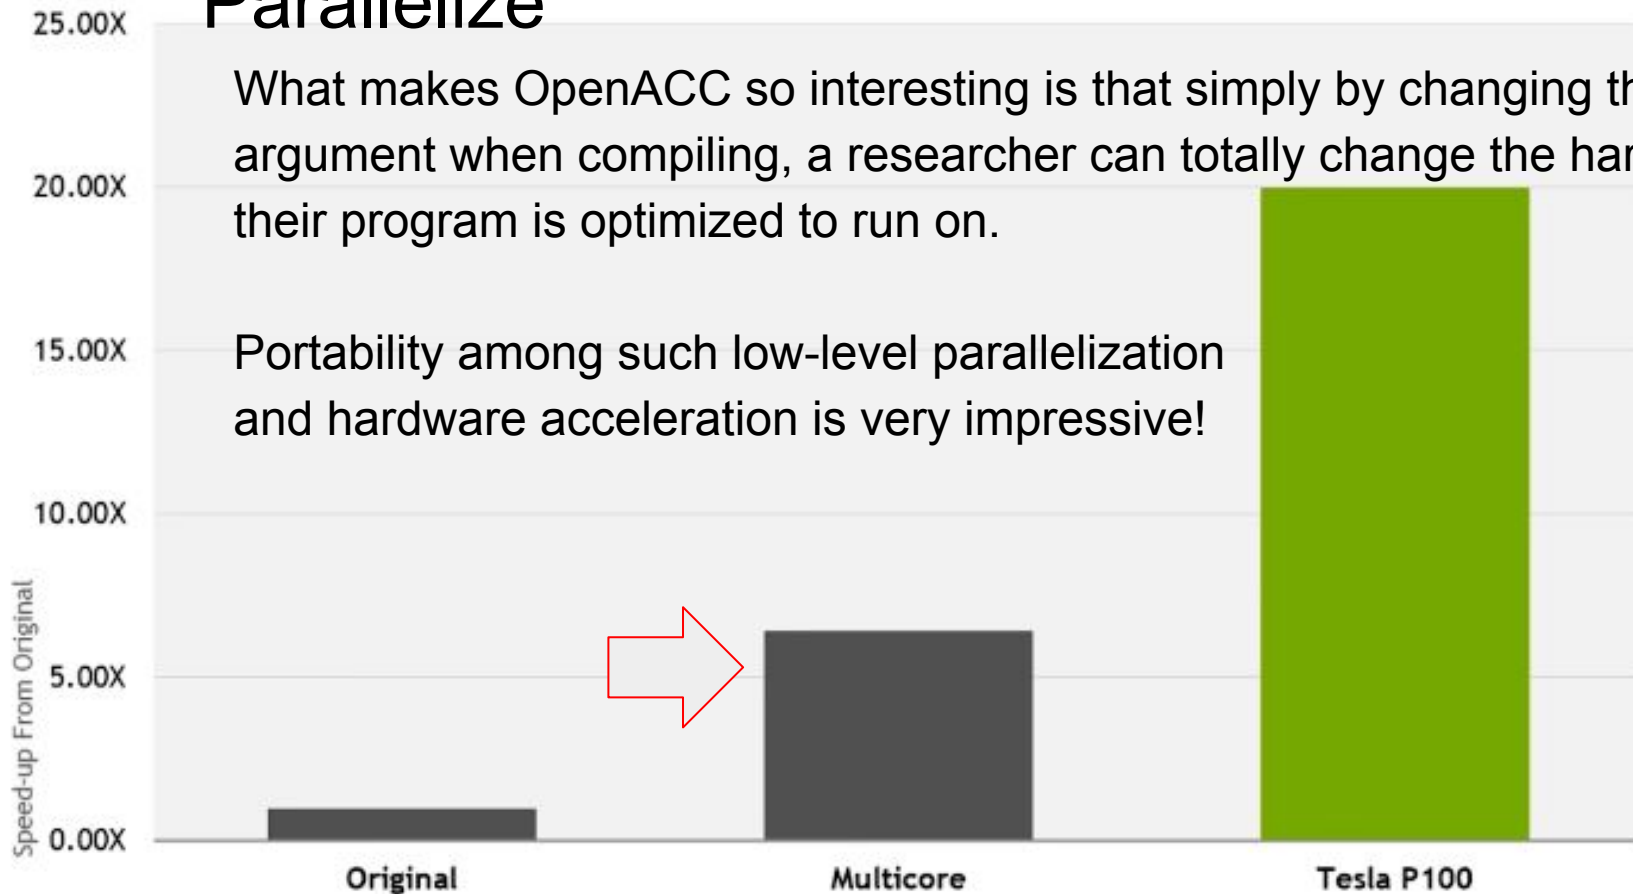
Without knowing anything about low-level hardware acceleration or parallelization, a researcher can use OpenACC directives to fully utilize their specialized processors.



Parallelize

What makes OpenACC so interesting is that simply by changing the -ta argument when compiling, a researcher can totally change the hardware their program is optimized to run on.

Portability among such low-level parallelization and hardware acceleration is very impressive!



Parallelize

Inserting pragma directives according to the data given by a profiler can make a program ready to be ported with acceleration across a lot of hardware.

As in Assignment 1 Part 2, it would normally take a bit of research and tricky low-level implementation to get this kind of acceleration on a heterogeneous computing system.

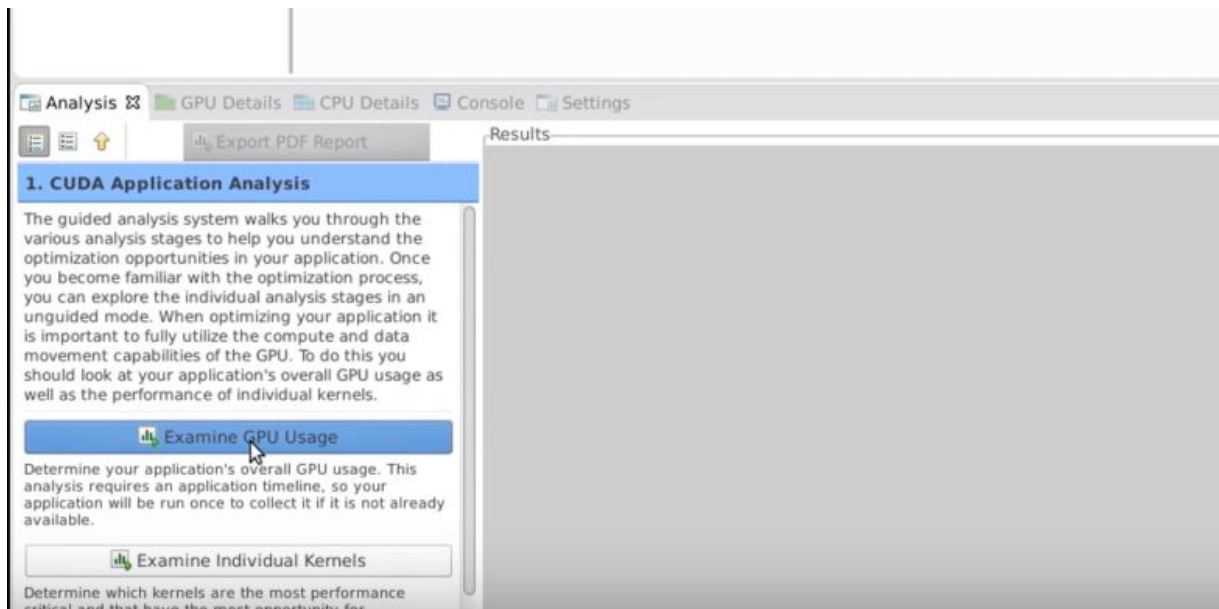
With OpenACC, simple directives allow you to achieve that acceleration across many sets of hardware easily.

Optimize

Optimizing with OpenACC means taking another look at the profiler and improving upon the compiler's best effort to parallelize automatically.

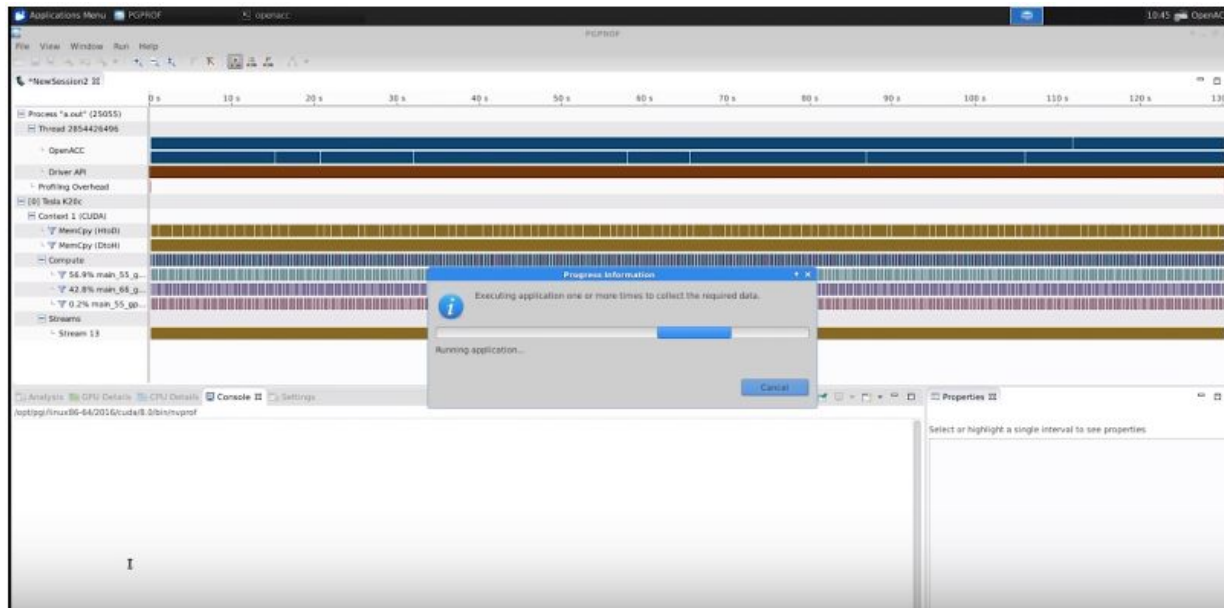
PGPROF will be used as an example.

Optimize



Here, the profiler will examine GPU usage.

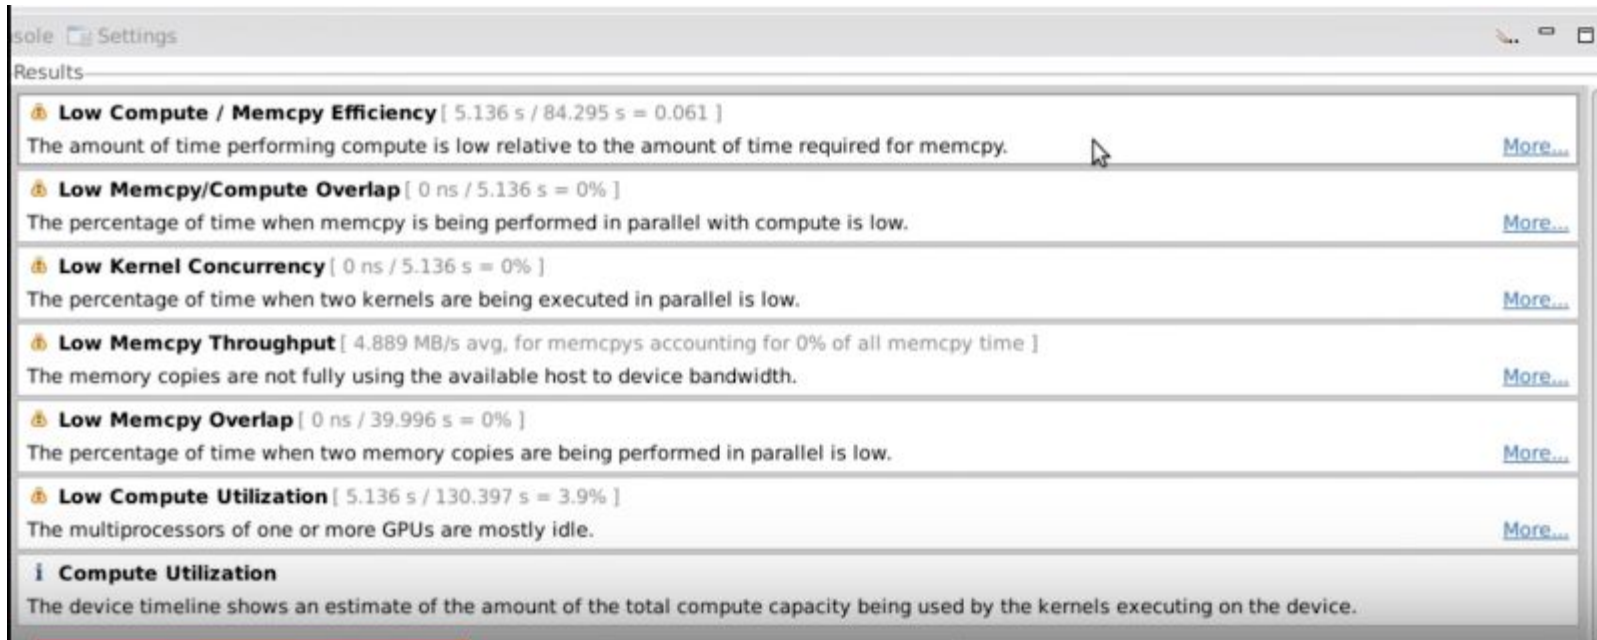
Optimize



“The guided analysis tool will rerun your code several times to help you determine what is limiting the performance of your application.”

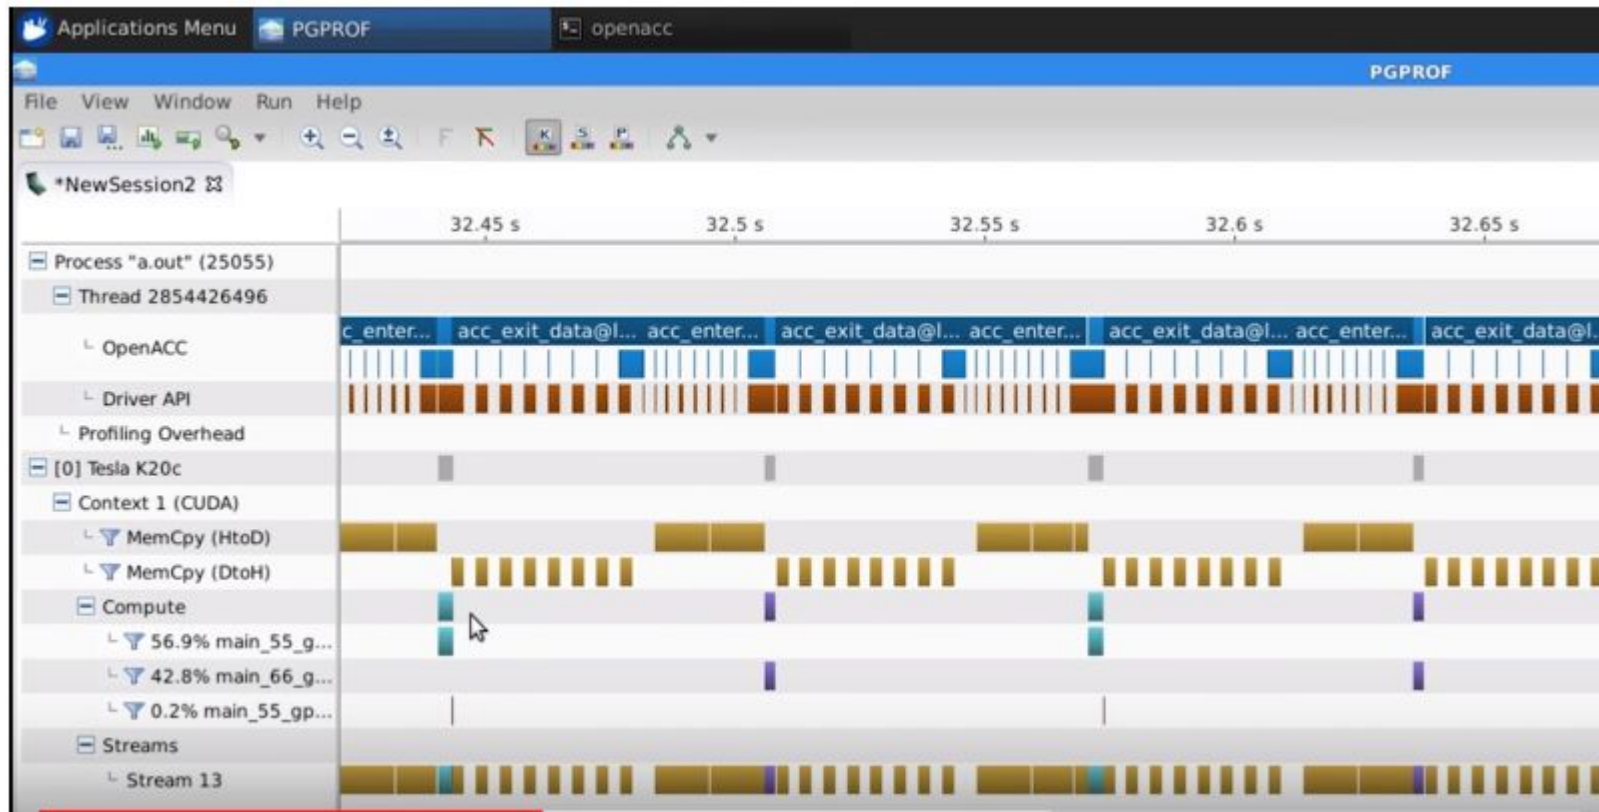
Source: “Step 3: Optimize” (1:04)

Optimize



The profiler found that data copies between the host and the accelerator were slowing down the code a lot. (Remember how we fixed this with AVX?)

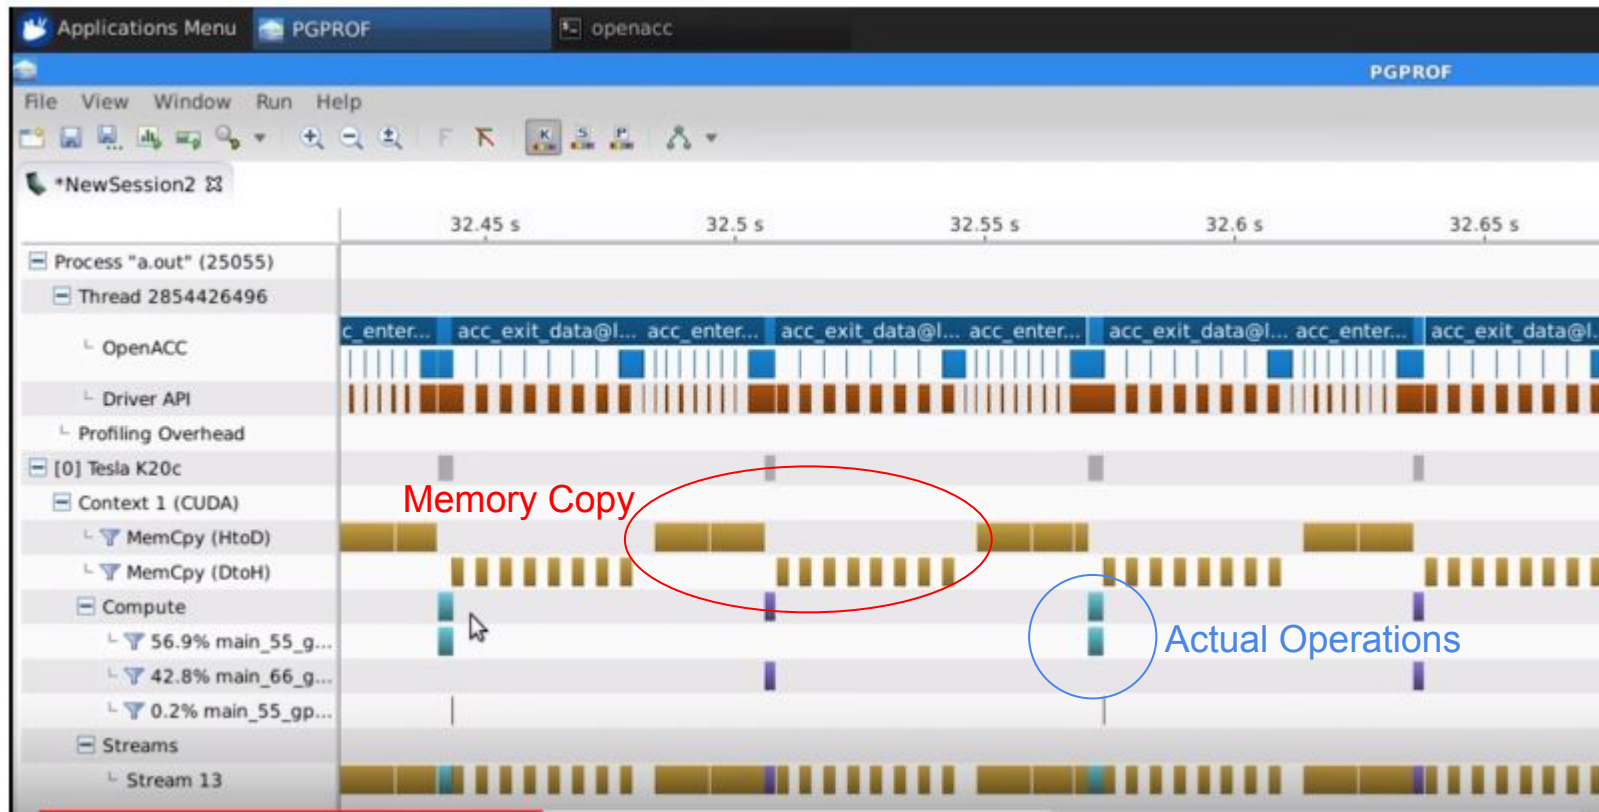
Optimize



The profiler even visualizes the execution time in charts, differentiating different accelerator operations as slices of time.

Source: "Step 3: Optimize"

Optimize



Note the very long MemCpy times compared to the Compute times!

Source: "Step 3: Optimize"

Optimize

```
#pragma acc data copy(A) create(Anew)
while ( error > tol && iter < iter_max )
{
    error = 0.0;

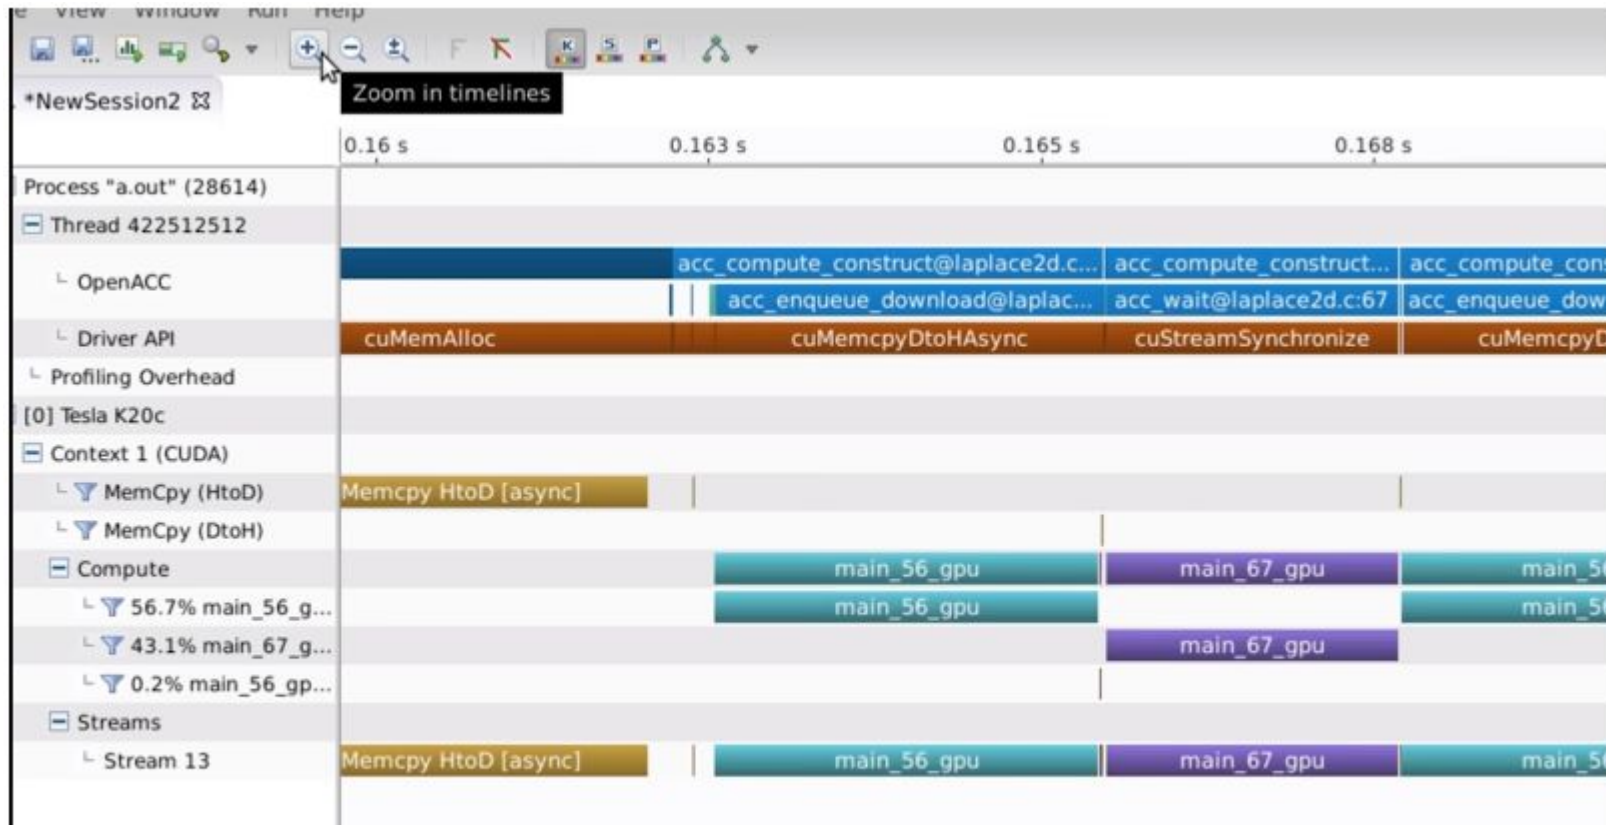
#pragma acc parallel loop reduction(max:error)
    for( int j = 1; j < n-1; j++)
    {
        for( int i = 1; i < m-1; i++ )
        {
            Anew[j][i] = 0.25 * ( A[j][i+1] + A[j][i-1]
                                + A[j-1][i] + A[j+1][i]);
            error = fmax( error, fabs(Anew[j][i] - A[j][i]));
        }
    }

#pragma acc parallel loop
    for( int j = 1; j < n-1; j++)
    {
        for( int i = 1; i < m-1; i++ )
        {
            A[i][i] = Anew[i][i];
        }
    }
}
```

Here, a data directive is used to manually tell the compiler how the data should get loaded to the accelerator (load input at start, unload result at end).

Source: “Step 3: Optimize” (2:18)

Optimize



That one directive “completely eliminated” the memory copy time! (Think about loadu_ps and load_ps).

Source: “Step 3: Optimize” (2:46)

Optimize

It's time to look at what OpenACC did wrong!

```
matvec(const matrix &, const vector &, const vector &):  
    8, include "matrix_functions.h"  
    12, Generating copyin(Acoefs[:A->nnz],cols[:A->nnz])  
        Generating implicit copyin(row_offsets[:num_rows+1])  
        Generating copyin(xcoefs[:num_rows])  
        Generating copyout(ycoefs[:num_rows])  
        Accelerator kernel generated  
        Generating Tesla code  
    16, #pragma acc loop gang /* blockIdx.x */  
    21, #pragma acc loop vector(128) /* threadIdx.x */  
        Generating reduction(+:sum)  
    21, Loop is parallelizable
```

OpenACC is parallelizing the loop with a vector length of 128, which means that it “operates on 128 loop iterations at one time.” (Think of our 256 bit AVX registers that operated 8 loop iterations at once.)

Source: “Step 3: Optimize” (2:59)

Optimize

```
1 #pragma once
2 #include "vector.h"
3 #include "matrix.h"
4
5 void matvec(const matrix& A, const vector& x, const vector &y)
6 {
7     unsigned int num_rows=A.num_rows;
8     unsigned int *restrict row_offsets=A.row_offsets;
9     unsigned int *restrict cols=A.cols;
10    double *restrict Acoefs=A.coefs;
11    double *restrict xcoefs=x.coefs;
12    double *restrict ycoefs=y.coefs;
13
14    #pragma acc parallel loop copyout(ycoefs[:num_rows]) \
15        copyin(Acoefs[:A.nnz],xcoefs[:num_rows],cols[:A.nnz])
16    for(int i=0;i<num_rows;i++) {
17        double sum=0;
18        int row_start=row_offsets[i];
19        int row_end=row_offsets[i+1];
20        #pragma acc loop reduction+:sum
21        for(int j=row_start;j<row_end;j++) {
22            unsigned int Acol=cols[j];
23            double Acoef=Acoefs[j];
24            double xcoef=xcoefs[Acol];
25            sum+=Acoef*xcoef;
26        }
27        ycoefs[i]=sum;
28    }
29 }
```

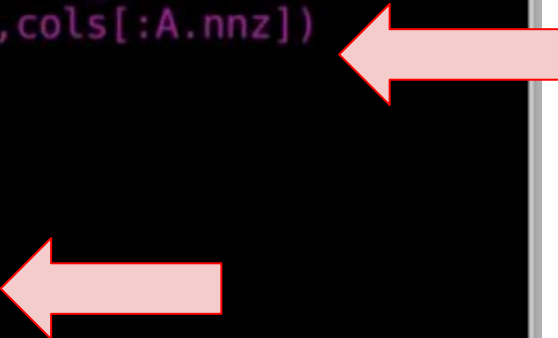
However, the loop does not even iterate 128 times, making an attempt to parallelize 128 times over pointless.



Optimize

These are the loops in question:

```
14 #pragma acc parallel loop copyout(ycoefs[:num_rows]) \
15     copyin(Acoefs[:A.nnz],xcoefs[:num_rows],cols[:A.nnz])
16     for(int i=0;i<num_rows;i++) {
17         double sum=0;
18         int row_start=row_offsets[i];
19         int row_end=row_offsets[i+1];
20 #pragma acc loop reduction(+:sum)
21         for(int j=row_start;j<row_end;j++) {
22             unsigned int Acol=cols[i];
```



Optimize

“I’ll reduce the vector length to 32, since this is the smallest vector length supported on my GPU.”

```
14 #pragma acc parallel loop copyout(ycoefs[:num_rows]) \
15     copyin(Acoefs[:A.nnz],xcoefs[:num_rows],cols[:A.nnz])
16     for(int i=0;i<num_rows;i++) {
17         double sum=0;
18         int row_start=row_offsets[i];
19         int row_end=row_offsets[i+1];
20 #pragma acc loop reduction(+:sum)
21         for(int j=row_start;j<row_end;j++) {
22             unsigned int Acol=cols[j];
```

Note the necessity for knowledge of low-level implementations at this point.

Optimize

“I’ll reduce the vector length to 32, since this is the smallest vector length supported on my GPU.”

```
14 #pragma acc parallel loop copyout(ycoefs[:num_rows]) \
15   copyin(Acoefs[:A.nnz],xcoefs[:num_rows],cols[:A.nnz])\
16   vector_length(32)
17   for(int i=0;i<num_rows;i++) {
18       double sum=0;
19       int row_start=row_offsets[i];
20       int row_end=row_offsets[i+1];
21 #pragma acc loop reduction(+:sum) vector
22   for(int j=row_start;j<row_end;j++) {
```

To fix what the compiler did wrong, we must explicitly vectorize the inner loop and specify a vector length of 32 in the outer loop. (Note that the outer loop directive specifies a “parallel region.”)

Optimize

“Because 32 is really too little parallelism for this GPU, I’ll tell the compiler to also parallelize the outer loop.”

```
14 #pragma acc parallel loop copyout(ycoefs[:num_rows]) \
15     copyin(Acoefs[:A.nnz],xcoefs[:num_rows],cols[:A.nnz])\
16     vector_length(32) worker num_workers(4) gang
17     for(int i=0;i<num_rows;i++) {
18         double sum=0;
19         int row_start=row_offsets[i];
20         int row_end=row_offsets[i+1];
21 #pragma acc loop reduction(+:sum) vector
22         for(int j=row_start;j<row_end;j++) {
23             unsigned int Acol=cols[i]:
```

This is done by assigning OpenACC “workers” and “gangs” to the loops, which further parallelizes the program.

Source: “Step 3: Optimize” (3:26)

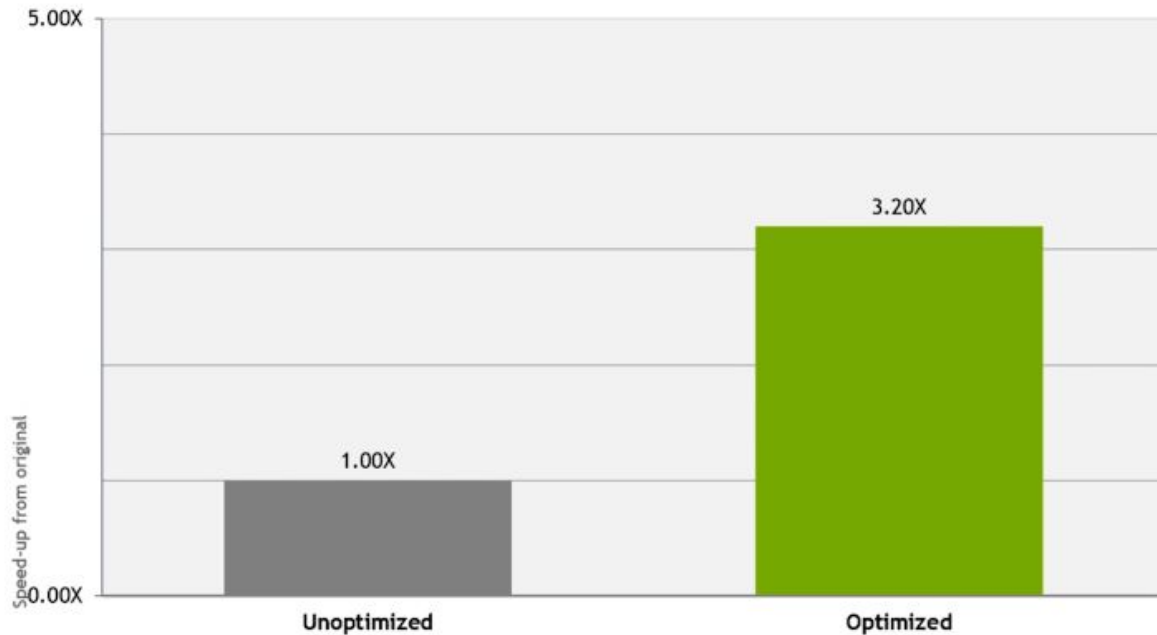
Optimize

Obviously, these optimizations require some low-level knowledge and are much more involved than just looking at the profiler output to find where to put directives.

“There’s no single optimization technique that will work on every code” (“Step 3: Optimize” 5:02).

Portability aside, some argue that it is more than worth the trouble to just learn how to use OpenCL instead, because the compiler is quite bad at automatically generating parallelization much of the time.

Optimize



Source: PGI 16.10, NVIDIA Tesla K100

“I roughly tripled the performance of my code by better mapping my loops to the hardware.”

Source: “Step 3: Optimize” (4:48)

Optimize

“I roughly tripled the performance of my code by better mapping my loops to the hardware” (“Step 3: Optimize” 4:48).

True, but isn't the whole point of OpenACC that your code is portable AND optimized between hardware?

At least OpenACC can handle everything under one set of directives though, and will always be portable across hardware.

Assignment 1: OpenACC Edition

We know firsthand how challenging it can be to write parallelized/vectorized code for a specific set of hardware.

Let's see how easy it is with OpenACC.

(And whether it works as well as our solutions)

Analyze

Analysis GPU Details (Summary) CPU Details ⌵ OpenACC Details Console Settings		
TOTAL ▼ Use the buttons on the top-right of this view to select how to display profile data		
Event	%	Time
▶ random	67.647%	2,070.01357 ms
▶ random_r	15.686%	480.00315 ms
main	10.784%	330.00216 ms
▶ rand	2.941%	90.00059 ms
▶ matvec(float const *, float const *)	1.961%	60.00039 ms
▶ ???	0.98%	30.0002 ms

Analyze

```
// Sequential implementation of a matrix vectorization
void matvec(const float *A, const float *x, unsigned int m, unsigned int n, float *b) {
    float vecSum;
    unsigned int in;
    for (unsigned int i = 0; i < m; i++) {
        vecSum = 0; // Reset sum to 0 after every row
        in = i*n;
        for (unsigned int j = 0; j < n; j++) {
            // Sum each product of corresponding matrix and vector values
            vecSum += (A[in + j] * x[j]);
        }
        // Record the sum in the resulting vector
        b[i] = vecSum;
    }
}
```

Analyze

```
// Sequential implementation of a matrix vectorization
void matvec(const float *A, const float *x, unsigned int m, unsigned int n, float *b) {
    float vecSum;
    unsigned int in;
    for (unsigned int i = 0; i < m; i++) {
        vecSum = 0; // Reset sum to 0 after every row
        in = i*n;
        for (unsigned int j = 0; j < n; j++) {
            // Sum each product of corresponding
            vecSum += (A[in + j] * x[j]);
        }
        // Record the sum in the resulting vector
        b[i] = vecSum;
    }
}
```

Multiple markers at this line

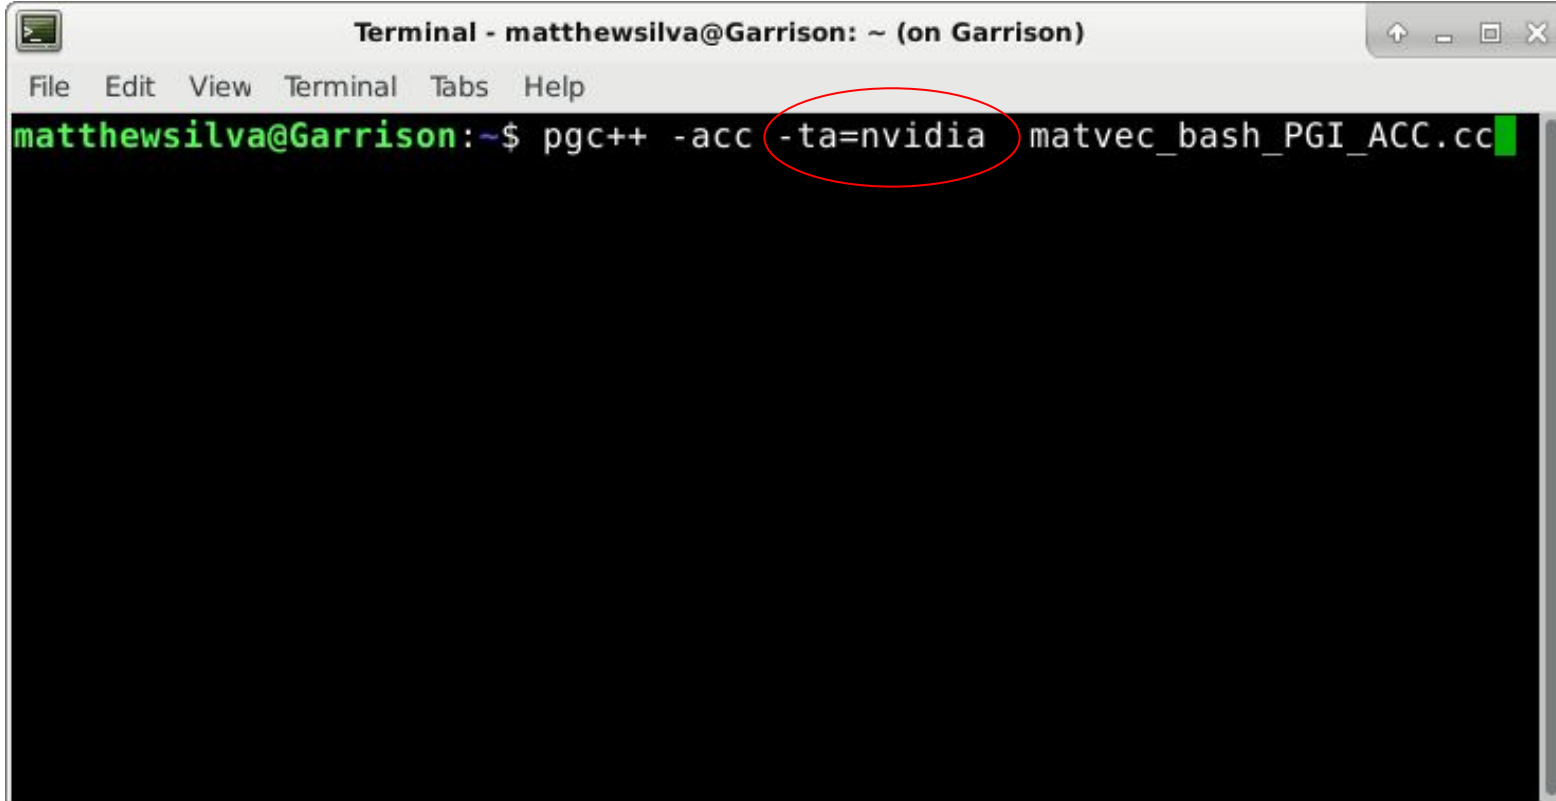
- FMA (fused multiply-add) instruction(s) generated
- Generated 2 prefetch instructions for the loop
- Generated vector simd code for the loop containing reductions

Parallelize

```
GNU nano 2.7.4 File: matvec bash PGI ACC.cc

// Sequential implementation of a matrix vectorization
void matvec(const float *A, const float *x, unsigned int m, unsigned int n, float *b) {
    float vecSum;
    unsigned int in;
    int mn = m*n;
    for (unsigned int i = 0; i < m; i++) {
        vecSum = 0; // Reset sum to 0 after every row
        in = i*n;
        #pragma acc parallel loop
        for (unsigned int j = 0; j < n; j++) {
            // Sum each product of corresponding matrix and vector values
            vecSum += (A[in + j] * x[j]);
        }
        // Record the sum in the resulting vector
        b[i] = vecSum;
    }
}
```

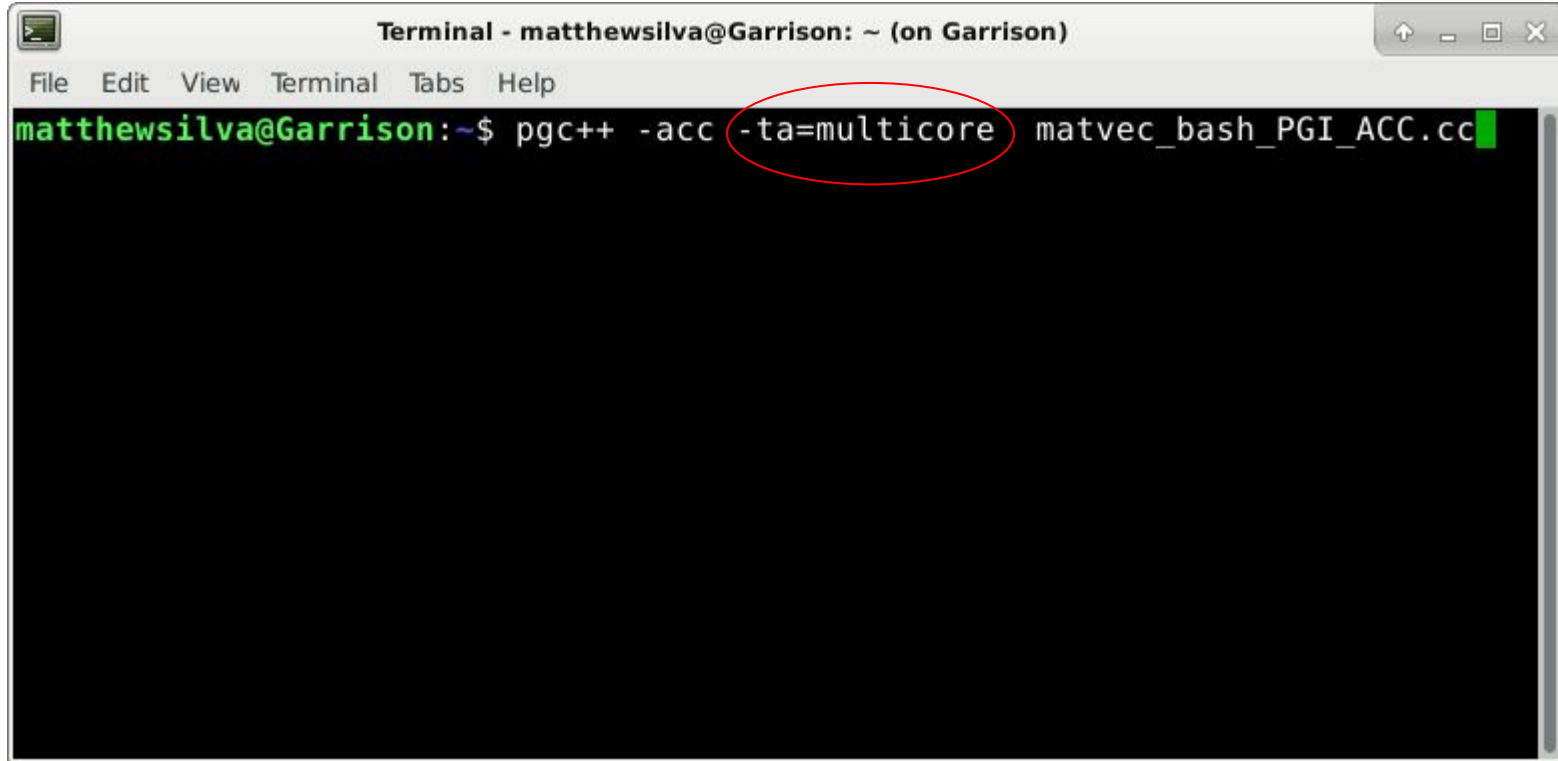
Parallelize (For Nvidia)



A terminal window titled "Terminal - matthewsilva@Garrison: ~ (on Garrison)". The window has a menu bar with "File", "Edit", "View", "Terminal", "Tabs", and "Help". The command prompt shows the user "matthewsilva@Garrison" at the shell "\$". The command being entered is "pgc++ -acc -ta=nvidia matvec_bash_PGI_ACC.cc". The option "-ta=nvidia" is circled in red. The terminal background is black, and the text is green.

```
matthewsilva@Garrison:~$ pgc++ -acc -ta=nvidia matvec_bash_PGI_ACC.cc
```


Parallelize (for Multicore CPU)




A terminal window titled "Terminal - matthewsilva@Garrison: ~ (on Garrison)" is shown. The window has a menu bar with "File", "Edit", "View", "Terminal", "Tabs", and "Help". The command prompt shows the user "matthewsilva@Garrison" at the shell "\$". The command being entered is "pgc++ -acc -ta=multicore matvec_bash_PGI_ACC.cc". The option "-ta=multicore" is circled in red. The terminal background is black, and the text is green. A green cursor is at the end of the command.

```
matthewsilva@Garrison:~$ pgc++ -acc -ta=multicore matvec_bash_PGI_ACC.cc
```

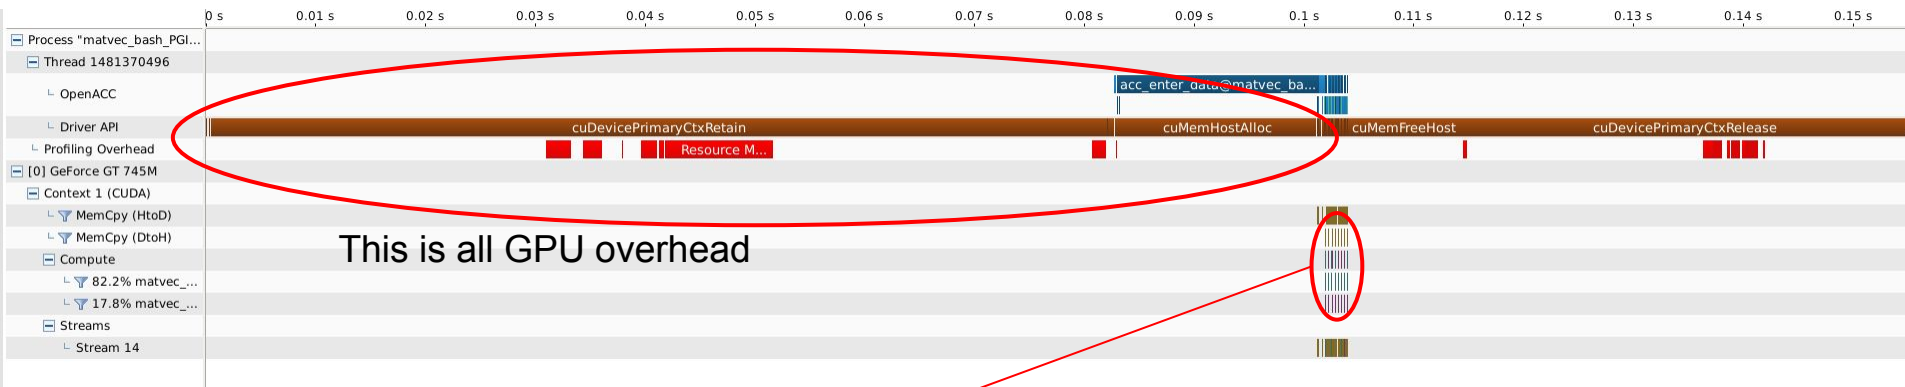

Parallelize



Parallelize

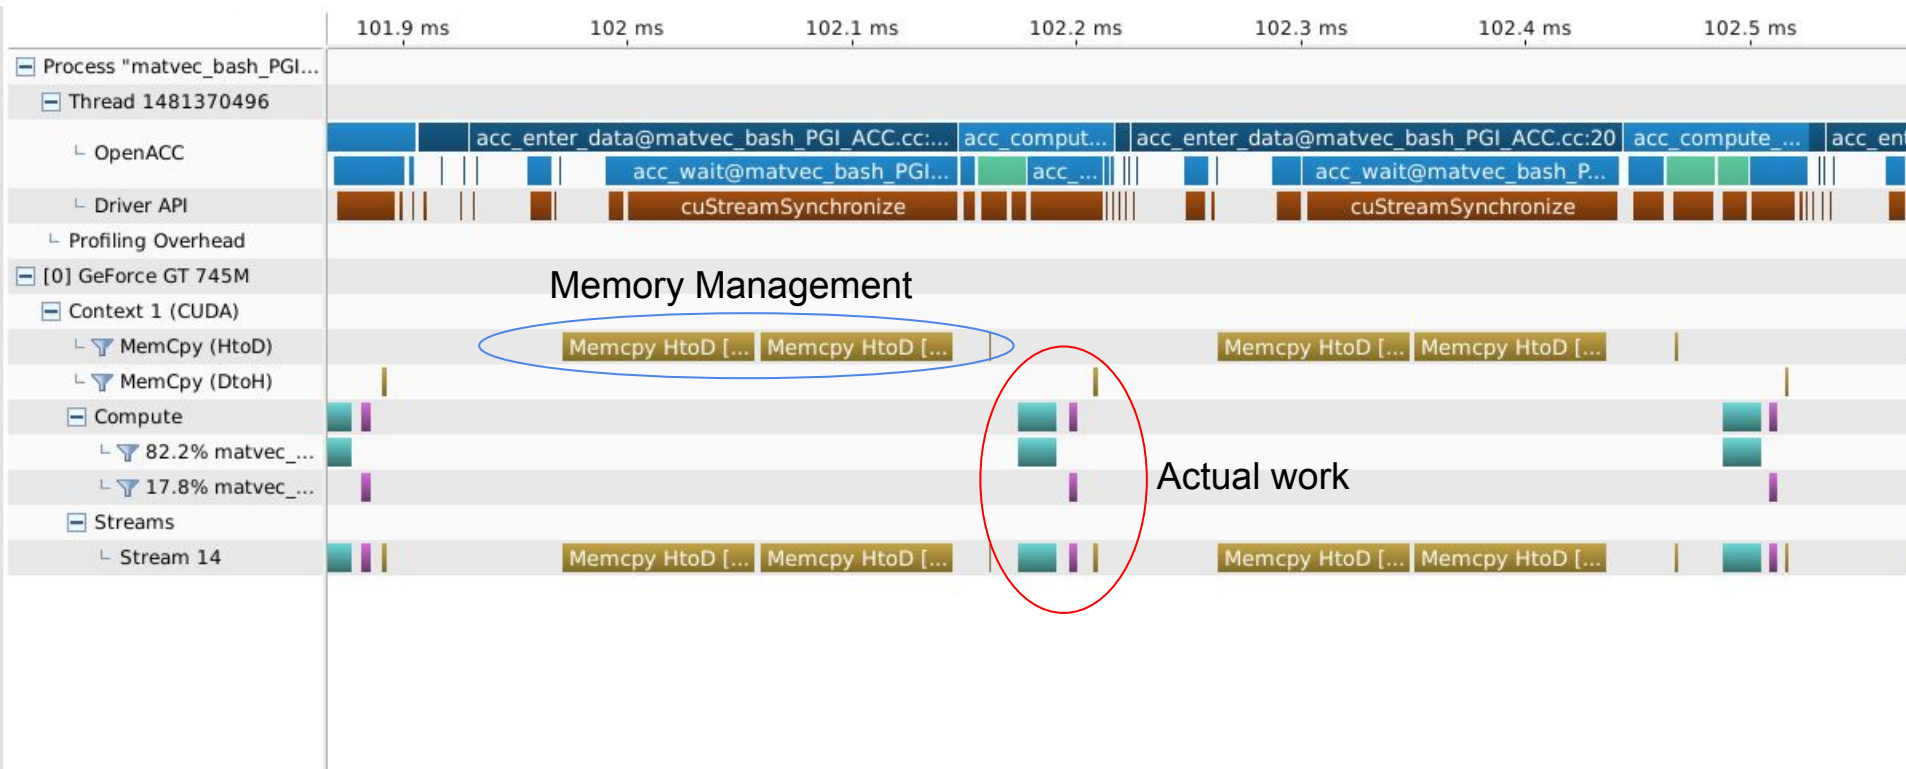


Optimize



This is where the memory management and work actually get done

Optimize



Optimize

Results

 **Low Compute / Malloc Efficiency** [170.183 μ s / 1.38581 ms = 0.123]

The amount of time performing compute is low relative to the amount of time required for malloc.

[More...](#)

 **Low Malloc/Compute Overlap** [0 ns / 170.183 μ s = 0%]

The percentage of time when malloc is being performed in parallel with compute is low.

[More...](#)

 **Low Kernel Concurrency** [0 ns / 170.183 μ s = 0%]

The percentage of time when two kernels are being executed in parallel is low.

[More...](#)

 **Low Malloc Throughput** [1.513 GB/s avg, for mallocs accounting for 100% of all malloc time]

The memory copies are not fully using the available host to device bandwidth.

[More...](#)

 **Low Compute Utilization** [170.183 μ s / 155.17625 ms = 0.1%]

The multiprocessors of one or more GPUs are mostly idle.

[More...](#)

Absolutely terrible!

Compute Utilization

The device timeline shows an estimate of the amount of the total compute capacity being used by the kernels executing on the device.

Optimize


GNU nano 2.7.4

File: matvec bash PGI ACC red copy.cc

```
// Sequential implementation of a matrix vectorization
void matvec(const float *A, const float *x, unsigned int m, unsigned int n, float *$
    float vecSum;
    unsigned int in;
    int mn = m*n;
    #pragma acc data copy(A[0:mn]) copy(x[0:n])
    for (unsigned int i = 0; i < m; i++) {
        vecSum = 0; // Reset sum to 0 after every row
        in = i*n;
        #pragma acc parallel loop reduction(+:vecSum)
        for (unsigned int j = 0; j < n; j++) {
            // Sum each product of corresponding matrix and vector values
            vecSum += (A[in + j] * x[j]);
        }
        // Record the sum in the resulting vector
        b[i] = vecSum;
    }
}
```

Optimize

Results

 **Low Compute / Malloc Efficiency** [48.75105 ms / 742.64716 ms = 0.066]

The amount of time performing compute is low relative to the amount of time required for malloc.

[More...](#)

 **Low Malloc/Compute Overlap** [0 ns / 48.75105 ms = 0%]

The percentage of time when malloc is being performed in parallel with compute is low.

[More...](#)

 **Low Kernel Concurrency** [0 ns / 48.75105 ms = 0%]

The percentage of time when two kernels are being executed in parallel is low.

[More...](#)

 **Low Malloc Throughput** [1.627 GB/s avg, for mallocs accounting for 100% of all malloc time]

The memory copies are not fully using the available host to device bandwidth.

[More...](#)

 **Low Compute Utilization** [48.75105 ms / 1.01286 s = 4.8%] **Still terrible but 48 times better!**

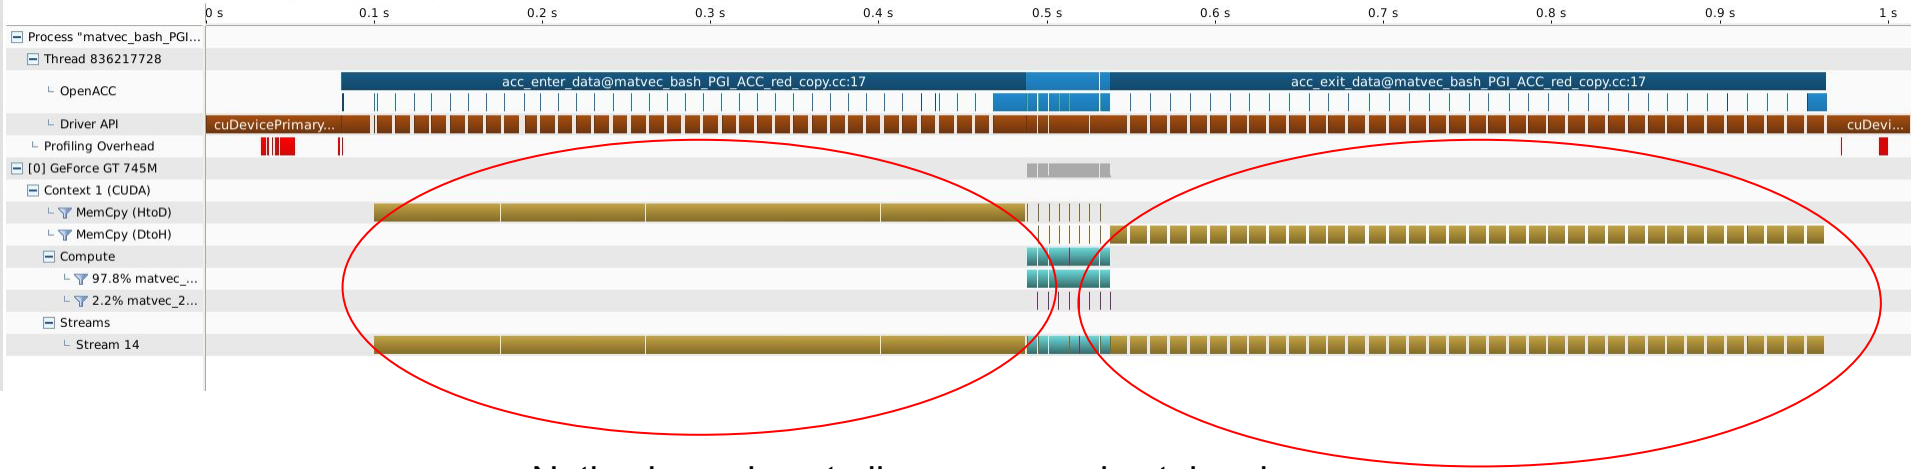
The multiprocessors of one or more GPUs are mostly idle.

[More...](#)

Compute Utilization

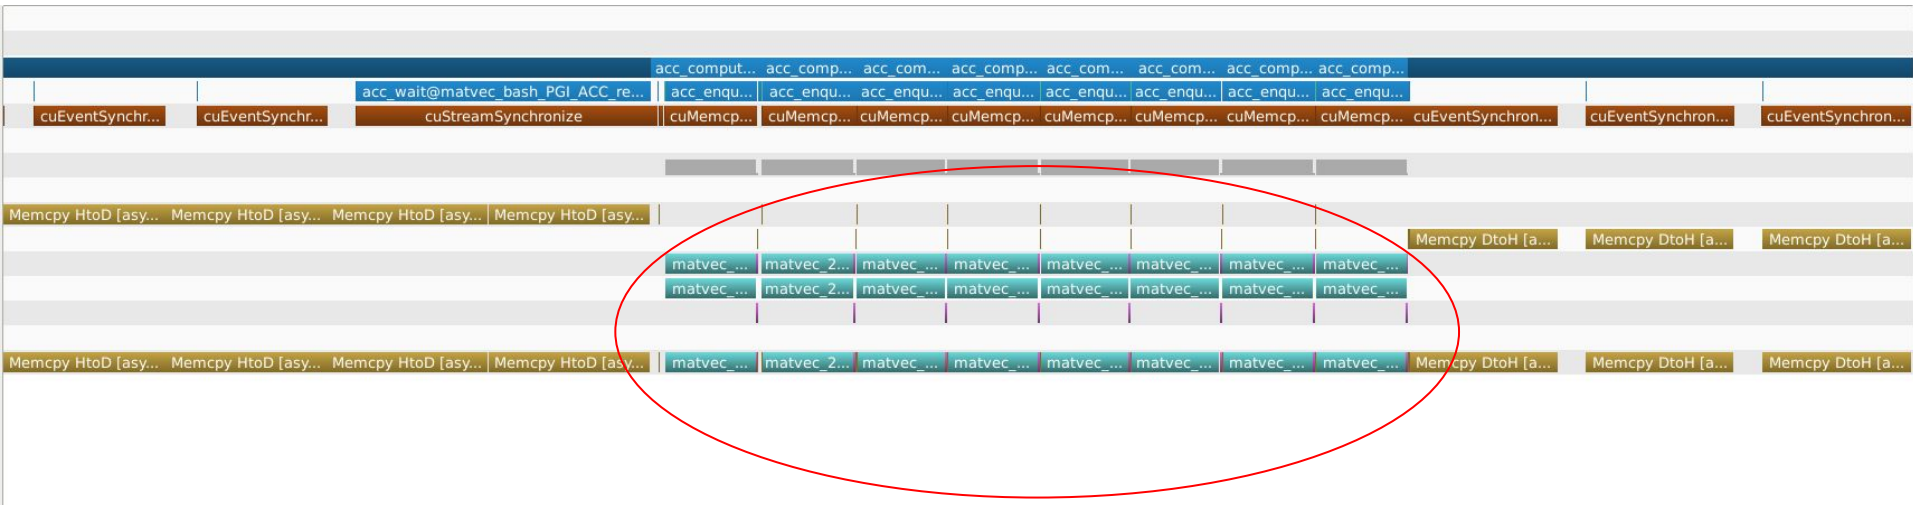
The device timeline shows an estimate of the amount of the total compute capacity being used by the kernels executing on the device.

Optimize



Notice how almost all memory copies take place before and after the computations


Optimize




All uninterrupted computations!

Optimize

OpenACC Optimization Speedups



Optimize



OpenACC Overview

- OpenACC: Compiler directives for automatically parallelizing code
- Performance-portable: Can optimize for different hardware (heterogeneous) automatically based on the same directives
- Profilers: Automatically point to parallelizable sections of code to guide directives
- Abstraction: More high-level than most attempts to accelerate code using specialized hardware
- Drawbacks: Compilers do not always follow the directives intelligently with respect to the hardware available

References

“Downloads & Tools.” *OpenACC*, OpenACC-standard.org, 2017.

“Get Started.” *OpenACC*, OpenACC-standard.org, 2017.

Melonakos. “No Free Lunch for GPU Compiler Directives.” *GPU and Accelerator Software Blog RSS*, GPU and Accelerator Software Blog, 11 Apr. 2012.

“The OpenACC API Quick Reference Guide.” *OpenACC*, OpenACC-standard.org, 2011.

“Research & Innovation.” *Softeco Sismat*, Softeco Sismat InformationTechnology, 2018.

“Step 1: Analyze, Three Steps to More Science Guide.” *NVIDIADeveloper*, NVIDIA, 13 Jan. 2017.

“Step 2: Parallelize, Three Steps to More Science Guide.” *NVIDIADeveloper*, NVIDIA, 13 Jan. 2017.

“Step 3: Optimize, Three Steps to More Science Guide.” *NVIDIADeveloper*, NVIDIA, 13 Jan. 2017.

“What Is OpenACC?” *OpenACC*, OpenACC-standard.org, 2017.